ECE 521 Assignment 2

Work Breakdown

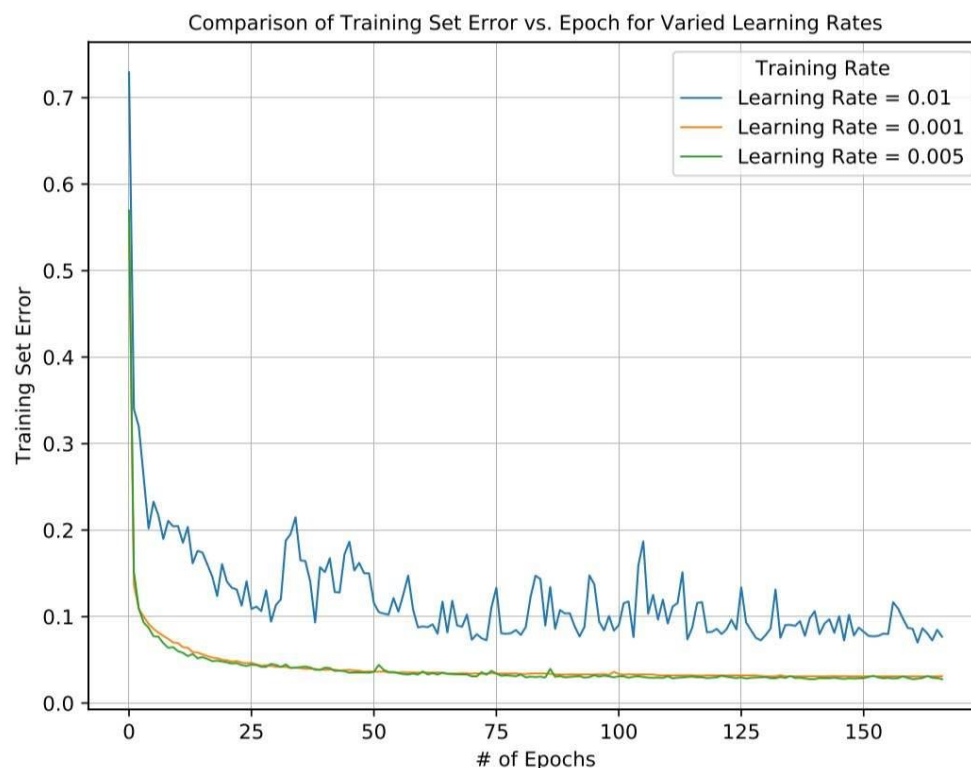| Group Member Name	Contribution Percentage
Jixong Deng	33%
Jeffrey Kirman	33%
Connor Smith	33%

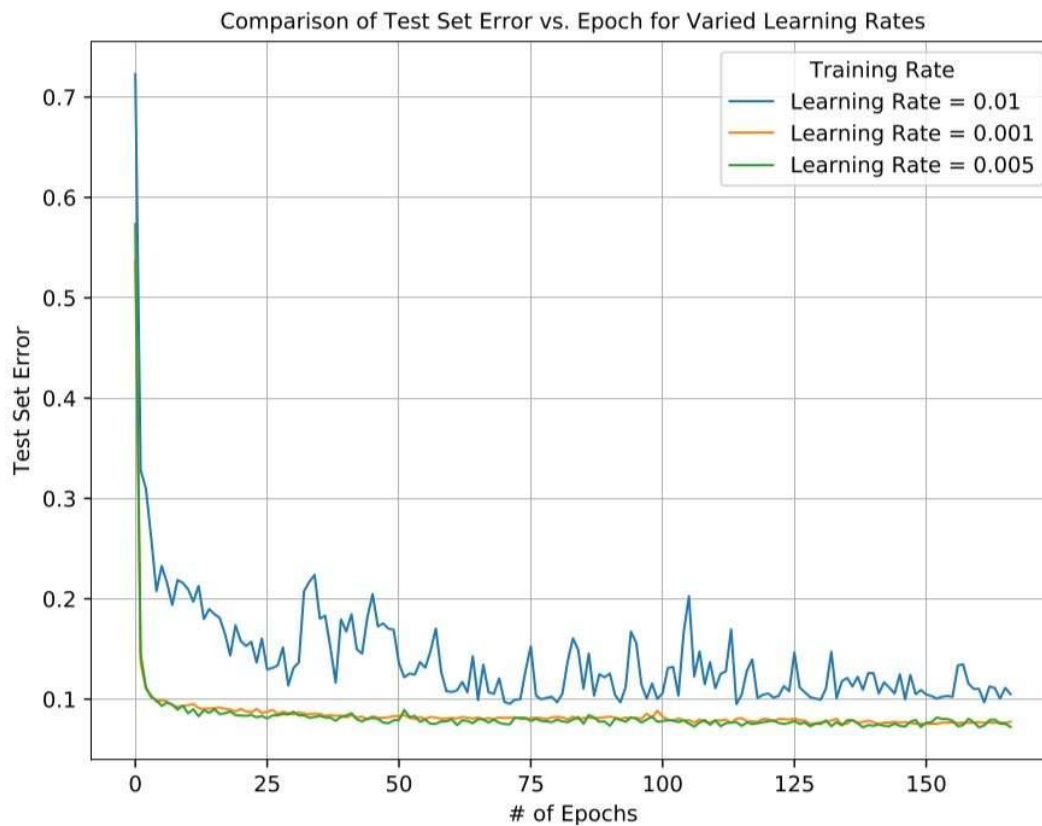## 1.1 – Feedforward fully connected neural networks

Part 1 – Layer-wise building block

The *create_new_layer()* function creates a new neural network layer from the input tensor *input_tensor* with a specified number of hidden units. The input tensor represents the raw X input vector without the bias padding. The weight matrix is initialized using Xavier initialization, and the bias term is zero initialized. The code for this layer creation function is available in **Append B – neuralnetworks.py**.
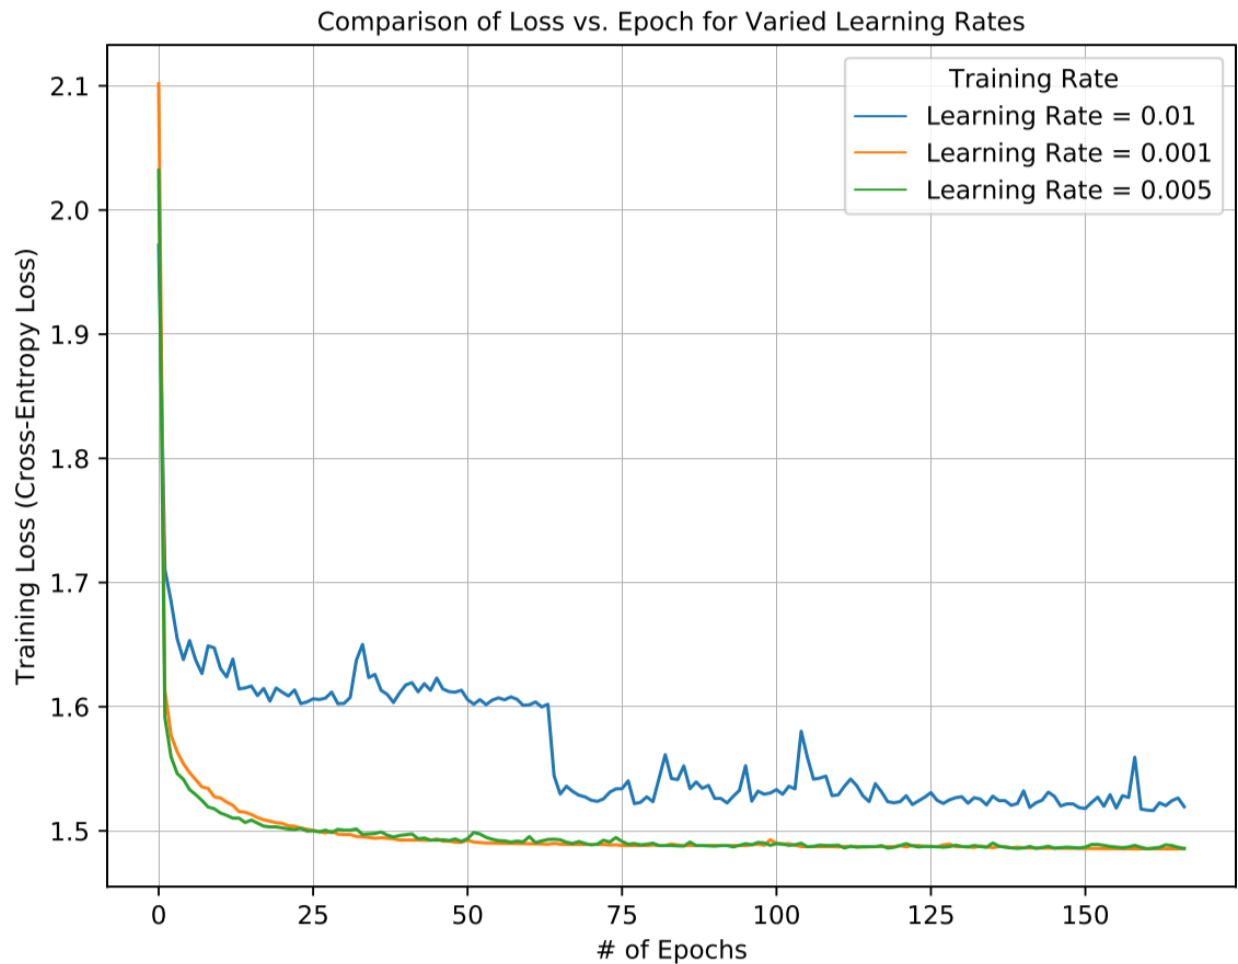
Part 2 – Learning

Using the above function, a simple one-layer neural network was created. To perform training, we used a minibatch size of 500 and 5000 iterations, totalling at 166 epochs. By setting the weight-decay coefficient $\lambda = 3 * 10^{-4}$, the following training, validation and test accuracy vs. epoch curves for different values of the learning rate η were observed:

Comparison of Validation Set Error vs. Epoch for Varied Learning Rates

Comparison of Test Set Error vs. Epoch for Varied Learning Rates

Plotting the training set loss vs. epoch, the following curve was observed:

Part 3 – Early stopping

Using the loss graph above and a learning rate of $\eta = 0.005$, we identify epoch 45 as the optimal early stopping point where loss did not appreciably change in later epochs. The test, training and validation accuracies at this point are recorded in the table below.

| Training Set Error	Validation Set Classification Error	Test Set Classification Error
0.0352	0.075	0.0811

1.2 – Effect of hyperparameters

Part 1 – Number of hidden units

The code used to see the difference in performance for different number of hidden units can be seen in **Appendix B - hyperparameters.py** specifically in the function *number_of_hidden_units()*. No weight decay was used in this model. The plots showing the training error and accuracies vs iterations (per epoch) can be seen in Appendix A. The following table summarizes the final values (best results bolded) for tuning the hyperparameters (where classification error $= 1 -$ accuracy):

# of nodes	100			500			1000		
η	0.01	0.005	**0.001**	0.01	0.005	**0.001**	0.01	0.005	**0.001**
Val. error	0.0800	0.0790	**0.0710**	0.0701	0.0670	**0.0670**	0.0760	0.0710	**0.0560**
Test error	0.0932	0.0870	**0.0896**	0.0808	0.0775	**0.0756**	0.0914	0.0793	**0.0727**

In summary, increasing the number of hidden units slightly reduces the validation error (by ~1%) for a single layer neural network for this dataset, but a lower number of nodes reduces computation time significantly.

Part 2 – Number of layers

2 Hidden Layers of 500 Nodes
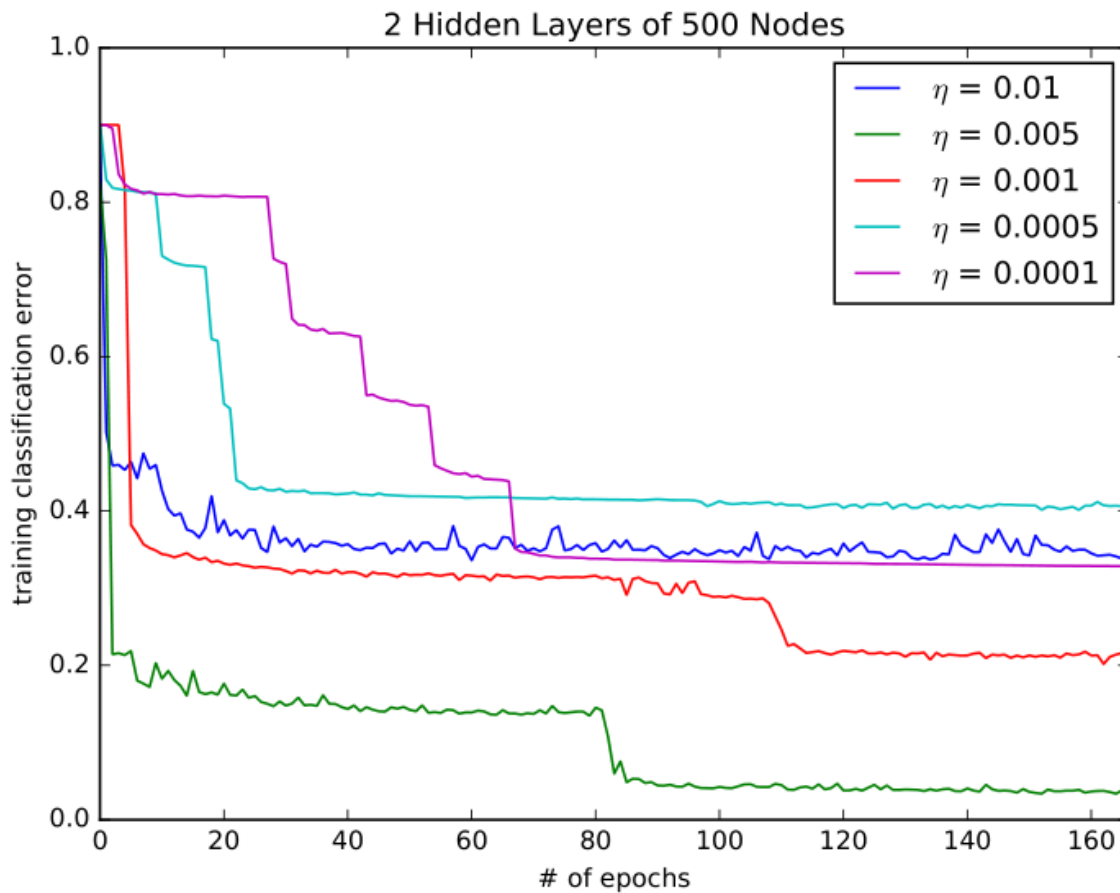
We can see in the graphs steep steps in decrease of error taking place at different epochs indicative of a multi-layer neural network. The following table summarizes the validation and test classification error final values for all learning rates. No weight decay was used in this model.

η	0.01	0.005	0.001	0.0005	0.0001
Val. error	0.373	**0.0800**	0.270	0.445	0.354
Test error	0.358	**0.0830**	0.262	0.452	0.356

Comparing the results of the test set for $\eta = 0.005$ in the 2 layer neural network, we see that the error is on par, if not a bit worse than the single layer neural network.

1.3 – Regularization

Part 1 – Dropout

The code used to see the add dropout to the neural network can be seen in **Appendix B - regularisation.py**. The following are the resulting plots. As seen, adding dropout to the single layer neural network slightly increases test accuracy (i.e. lowers test classification error) but has a slower rate of convergence.

Classification Error without Dropout

Classification Error with Dropout
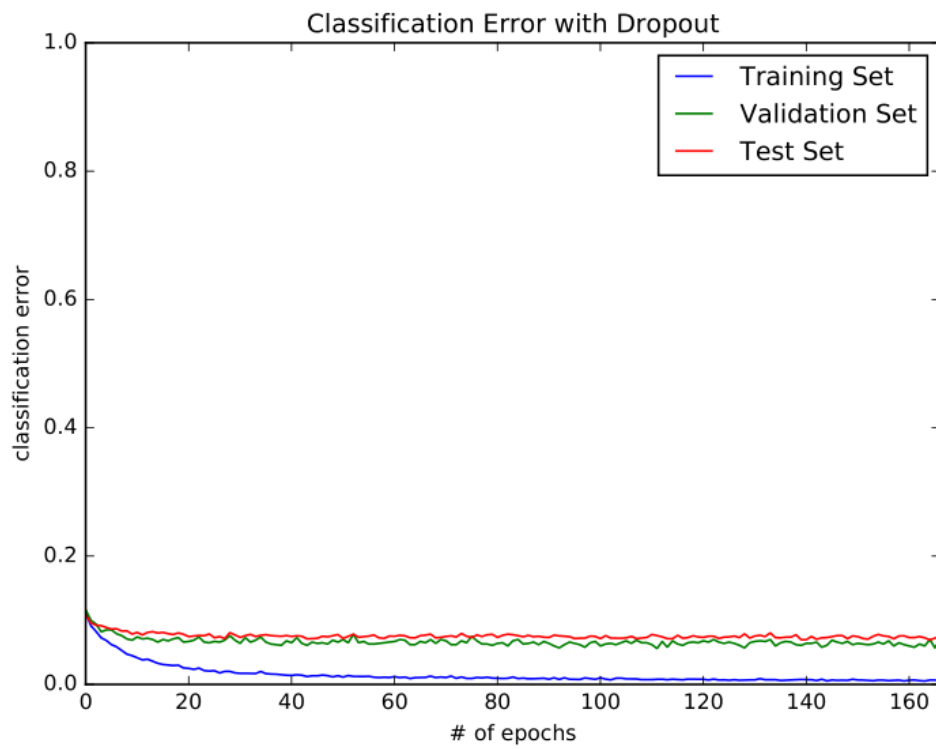
Part 2 – Visualization

100 of the neurons are displayed in the images before at 25% and 100% completion of training for non-dropout and dropout cases. We can see that from 25% to 100% the images start to become clearer and you can start to make out letters for some neurons. We also see that the neurons in to the dropout case are in general fuzzier than that of the non-dropout case.

25% complete without dropout

100% complete without dropout

25% complete with dropout

100% complete with dropout

1.4 – Exhaustive search

Part 1 – Random search

Randomization was added to the code for 1.3 which randomized the hyperparameters (the code can be seen in **Appendix B – exhaustive.py**). The following are the results after 166 epochs of training.

η	0.00149	0.00645	0.00724	0.00227	0.00531
# of nodes per layer	113	176	122	151	488
# of layers	5	1	4	2	2
Weight decay	7.799e-06	1.0217e-09	2.0157e-05	1.6428e-09	1.0331e-07
Dropout	True	True	False	False	False
Validation classification error	0.196	0.070	0.0820	0.0620	0.0650
Test classification error	0.195	0.081	0.0932	0.0761	0.0772

Part 2 – Exchanging ideas

We compared our results with several other groups in the class. Most other groups managed similar error values ranging from 7-15% on the test dataset, with the best test classification with its associated parameters summarized below (which bested even our results):

Learning Rate: 0.0016615

Layers: 2

Nodes per Layer: 458
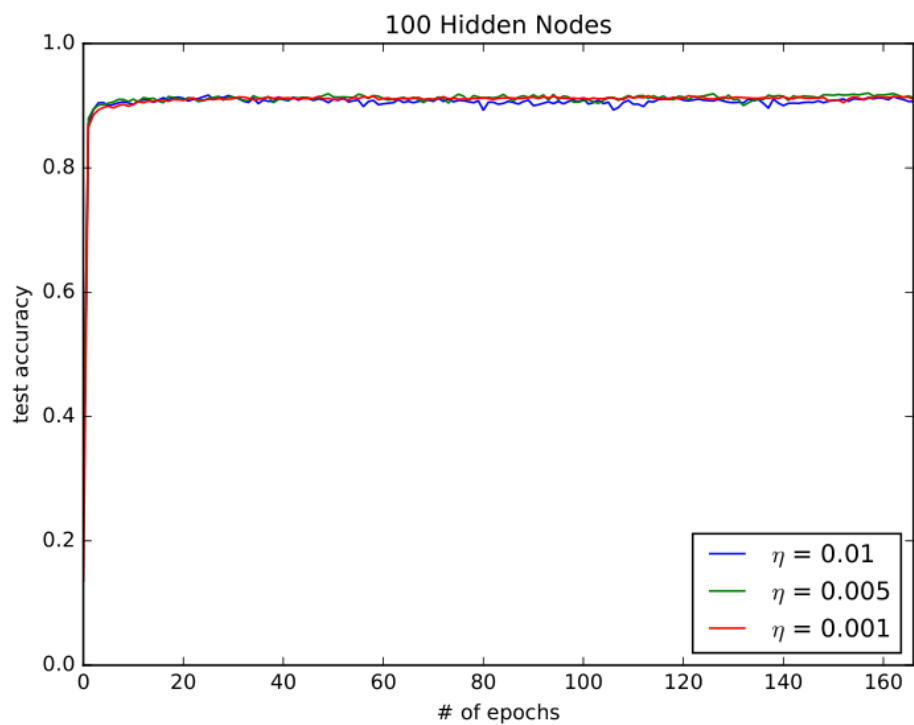
Dropout: Yes

Weight Decay Coefficient: 0.0002480

**Resultant Test Classification Error**: 0.0520

Appendix A – Graphs

1.2.1:

Jixong Deng (1001304741)
Jeffrey Kirman (1004387424)
Connor Smith (1000421411)

100 Hidden Nodes

100 Hidden Nodes

1000 Hidden Nodes

1000 Hidden Nodes

Appendix B – Python code

neural_networks.py

```python
import tensorflow as tf
import numpy as np


# Data loader for notMNIST dataset
def load_notmnist_data():
    with np.load("notMNIST.npz") as data:
        Data, Target = data["images"], data["labels"]
        np.random.seed(521)
        randIndx = np.arange(len(Data))
        np.random.shuffle(randIndx)
        Data = Data[randIndx]/255
        Target = Target[randIndx]
        trainData, trainTarget = Data[:15000], Target[:15000]
        t = np.zeros((trainTarget.shape[0], 10))
        t[np.arange(trainTarget.shape[0]), trainTarget] = 1
        trainTarget = t
        validData, validTarget = Data[15000:16000], Target[15000:16000]
        t = np.zeros((validTarget.shape[0], 10))
        t[np.arange(validTarget.shape[0]), validTarget] = 1
        validTarget = t
        testData, testTarget = Data[16000:], Target[16000:]
        t = np.zeros((testTarget.shape[0], 10))
        t[np.arange(testTarget.shape[0]), testTarget] = 1
        testTarget = t
        return (trainData.reshape(trainData.shape[0], -1), trainTarget, validData.reshape(validData.shape[0], -1),
validTarget, testData.reshape(testData.shape[0], -1), testTarget)


# Q1.1.1 layer-wise building block
def create_new_layer(input_tensor, num_hidden_units):
    '''
        @param input_tensor - outputs of the previous layer in the neural network, without the bias term.
        @param num_hidden_units - number of hidden units to use for this new layer
    '''
    # Create the new layer weight matrix using Xavier initialization
    input_dim = int(input_tensor.shape[-1])
    initializer = tf.contrib.layers.xavier_initializer()
    W_shape = [input_dim, num_hidden_units]
    W = tf.get_variable("Layer1_W", initializer=initializer(W_shape), dtype=tf.float32)
    # todo: zero initializer?
    b = tf.get_variable("Layer1_b", shape=[1, num_hidden_units], dtype=tf.float32)

    # MatMul the extended input tensor by the new weight matrix and add the biases
    output_tensor = tf.matmul(input_tensor, W) + b

    # Return this operation
    return output_tensor


# Q1.1.2 learning
def learning():
    xTrain, yTrain, xValid, yValid, xTest, yTest = load_notmnist_data()

    with tf.Graph().as_default():
        num_hidden_units = 1000
        decay = 0
        B = 500
        learning_rates = [0.01, 0.005, 0.001]
        iters = 5000

        num_iters_per_epoch = len(xTrain)//B # number of iterations we have to do for one epoch
        print("Num epochs = ",iters/num_iters_per_epoch)

        # hyperparameters
        learning_rate = tf.placeholder(dtype=tf.float32, name="learning-rate")

        # Get Data
        xTrainTensor = tf.constant(xTrain, dtype=tf.float32, name="X-Training")
        yTrainTensor = tf.constant(yTrain, dtype=tf.float32, name="Y-Training")
        xTestTensor = tf.constant(xTest, dtype=tf.float32, name="X-Test")
        yTestTensor = tf.constant(yTest, dtype=tf.float32, name="Y-Test")
```

```
        xValidTensor = tf.constant(xValid, dtype=tf.float32, name="X-Validation")
        yValidTensor = tf.constant(yValid, dtype=tf.float32, name="Y-Validation")

        Xslice, yslice = tf.train.slice_input_producer([xTrainTensor, yTrainTensor], num_epochs=None)

        Xbatch, ybatch = tf.train.batch([Xslice, yslice], batch_size = B)

        with tf.variable_scope("default") as scope:
            # Create neural network layers for training
            trainb_batchOutput = create_new_layer(Xbatch, num_hidden_units)
            trainb_activatedOutput = tf.nn.relu(trainb_batchOutput)

            scope.reuse_variables()
            layer1_w = tf.get_variable("Layer1_W", shape=[784, num_hidden_units], dtype=tf.float32)
            layer1_b = tf.get_variable("Layer1_b", shape=[1, num_hidden_units], dtype=tf.float32)

            train_output = tf.matmul(xTrainTensor, layer1_w) + layer1_b
            train_activatedOutput = tf.nn.relu(train_output)

            valid_output = tf.matmul(xValidTensor, layer1_w) + layer1_b
            valid_activatedOutput = tf.nn.relu(valid_output)

            test_output = tf.matmul(xTestTensor, layer1_w) + layer1_b
            test_activatedOutput = tf.nn.relu(test_output)

            outputWeights_size = [int(trainb_activatedOutput.shape[-1]), 10] # We want a [1,10] tensor to get
probabilities for each class
            outputWeights = tf.Variable(tf.contrib.layers.xavier_initializer()(outputWeights_size), name="Output_W")
            outputBias = tf.Variable(0, dtype=tf.float32, name="Output_Bias")

            trainb_y_pred = tf.sigmoid(tf.matmul(trainb_activatedOutput, outputWeights) + outputBias)
            train_y_pred = tf.sigmoid(tf.matmul(train_activatedOutput, outputWeights) + outputBias)
            valid_y_pred = tf.sigmoid(tf.matmul(valid_activatedOutput, outputWeights) + outputBias)
            test_y_pred = tf.sigmoid(tf.matmul(test_activatedOutput, outputWeights) + outputBias)
            trainb_softmaxLoss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=trainb_y_pred,
labels=ybatch)) + decay * tf.nn.l2_loss(layer1_w)
            train_softmaxLoss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=train_y_pred,
labels=yTrainTensor)) + decay * tf.nn.l2_loss(layer1_w)

            train_accuracy = tf.count_nonzero(tf.equal(tf.argmax(train_y_pred, 1), tf.argmax(yTrainTensor, 1))) /
yTrainTensor.shape[0]
            valid_accuracy = tf.count_nonzero(tf.equal(tf.argmax(valid_y_pred, 1), tf.argmax(yValidTensor, 1))) /
yValidTensor.shape[0]
            test_accuracy = tf.count_nonzero(tf.equal(tf.argmax(test_y_pred, 1), tf.argmax(yTestTensor, 1))) /
yTestTensor.shape[0]

        # optimizer function
        optimizer = tf.train.AdamOptimizer(learning_rate).minimize(trainb_softmaxLoss)

        # TODO run a lot of iterations, plot loss vs epochs and classification error vs epochs
        for r in learning_rates:
            loss_amounts = []
            train_accs = []
            test_accs = []
            valid_accs = []
            with tf.Session() as sess:
                coord = tf.train.Coordinator()
                threads = tf.train.start_queue_runners(sess=sess, coord=coord)
                sess.run(tf.global_variables_initializer())
                sess.run(tf.local_variables_initializer())
                for i in range(iters):
                    sess.run([optimizer], feed_dict={learning_rate: r})
                    if (i % num_iters_per_epoch == 0):
                        t_loss, t_acc, v_acc, test_acc = sess.run([train_softmaxLoss, train_accuracy, valid_accuracy,
test_accuracy])
                        print("Epoch: {}, Training Loss: {}, Accuracies: [{}, {}, {}]".format(i//num_iters_per_epoch,
t_loss, t_acc, v_acc, test_acc))
                        loss_amounts.append(t_loss)
                        train_accs.append(t_acc)
                        test_accs.append(test_acc)
                        valid_accs.append(v_acc)
                np.save("1.1.2_r{}_loss".format(r), loss_amounts)
                np.save("1.1.2_r{}_train_acc".format(r), train_accs)
                np.save("1.1.2_r{}_test_acc".format(r), test_accs)
                np.save("1.1.2_r{}_valid_acc".format(r), valid_accs)
learning()
```

hyperparameters.py

```python
import numpy as np
import tensorflow as tf

def load_notMNIST():
    with np.load("notMNIST.npz") as data:
        Data, Target = data["images"], data["labels"]
        np.random.seed(521)
        randIndx = np.arange(len(Data))
        np.random.shuffle(randIndx)
        Data = Data[randIndx]/255
        Target = Target[randIndx]
        trainData, trainTarget = Data[:15000], Target[:15000]
        t = np.zeros((trainTarget.shape[0], 10))
        t[np.arange(trainTarget.shape[0]), trainTarget] = 1
        trainTarget = t
        validData, validTarget = Data[15000:16000], Target[15000:16000]
        t = np.zeros((validTarget.shape[0], 10))
        t[np.arange(validTarget.shape[0]), validTarget] = 1
        validTarget = t
        testData, testTarget = Data[16000:], Target[16000:]
        t = np.zeros((testTarget.shape[0], 10))
        t[np.arange(testTarget.shape[0]), testTarget] = 1
        testTarget = t
        return (trainData.reshape(trainData.shape[0], -1), trainTarget, validData.reshape(validData.shape[0], -1),
validTarget, testData.reshape(testData.shape[0], -1), testTarget)

def create_new_layer(input_tensor, num_hidden_units):
    '''
        @param input_tensor - outputs of the previous layer in the neural network, without the bias term.
        @param num_hidden_units - number of hidden units to use for this new layer
    '''
    # Create the new layer weight matrix using Xavier initialization
    input_dim = int(input_tensor.shape[-1])
    initializer = tf.contrib.layers.xavier_initializer()
    W_shape = [input_dim, num_hidden_units]
    W = tf.get_variable("W", initializer=initializer(W_shape), dtype=tf.float32)
    # todo: zero initializer?
    b = tf.get_variable("b", shape=[1, num_hidden_units], dtype=tf.float32)

    # MatMul the extended input tensor by the new weight matrix and add the biases
    output_tensor = tf.matmul(input_tensor, W) + b

    # Return this operation
    return output_tensor

def number_of_hidden_units():
    # Constants
    B = 500
    iters = 5000
    learning_rates = [0.01, 0.005, 0.001]
    hidden_units = [100,500,1000]
    output_data = [[],[],[]]

    # Load data
    (trainData, trainTarget, validData, validTarget,
        testData, testTarget) = load_notMNIST()

    # Precalculations
    num_iters_per_epoch = len(trainData)//B # number of iterations we have to do for one epoch
    print("Num epochs = ",iters/num_iters_per_epoch)
    inds = np.arange(trainData.shape[0])

    # Set place-holders & variables
    X = tf.placeholder(tf.float32, shape=(None, trainData.shape[-1]), name='X')
    Y = tf.placeholder(tf.float32, shape=(None, 10), name='Y')
    learning_rate = tf.placeholder(tf.float32, name='learning-rate')

    for h in range(0, len(hidden_units)):
        for lr in range(len(learning_rates)):
            # Build graph
            with tf.variable_scope("layer1_"+str(hidden_units[h])+"_"+str(lr), reuse=tf.AUTO_REUSE):
                s_1 = create_new_layer(X, hidden_units[h])
            x_1 = tf.nn.relu(s_1)
            with tf.variable_scope("layer2_"+str(hidden_units[h])+"_"+str(lr), reuse=tf.AUTO_REUSE):
                s_2 = create_new_layer(x_1, 10)
```

```
            x_2 = tf.nn.softmax(s_2)

            # Calculate loss & accuracy
            loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=s_2, labels=Y))
            accuracy = tf.reduce_mean(tf.cast(tf.equal(tf.argmax(x_2, 1), tf.argmax(Y, 1)), tf.float32))

            print("Number of hidden units", hidden_units[h])

            with tf.Session() as sess:
                with tf.variable_scope("default", reuse=tf.AUTO_REUSE):
                    optimizer = tf.train.AdamOptimizer(learning_rate).minimize(loss)
                    coord = tf.train.Coordinator()
                    threads = tf.train.start_queue_runners(sess=sess, coord=coord)
                    sess.run(tf.global_variables_initializer())
                    sess.run(tf.local_variables_initializer())
                    print("Learning rate = ",learning_rates[lr])
                    temp_output = []
                    for i in range(iters):
                        if (i % num_iters_per_epoch == 0):
                            np.random.shuffle(inds)
                        sess.run([optimizer], feed_dict={learning_rate: learning_rates[lr],
                                X: trainData[inds[B*(i%num_iters_per_epoch):B*((i+1)%num_iters_per_epoch)]],
                                Y: trainTarget[inds[B*(i%num_iters_per_epoch):B*((i+1)%num_iters_per_epoch)]]})
                        if (i % num_iters_per_epoch == 0):
                            t_loss, t_acc = sess.run([loss, accuracy], feed_dict={X: trainData, Y: trainTarget})
                            v_loss, v_acc = sess.run([loss, accuracy], feed_dict={X: validData, Y: validTarget})
                            test_loss, test_acc = sess.run([loss, accuracy], feed_dict={X: testData, Y: testTarget})
                            print("Epoch: {}, Training Loss: {}, Accuracies: [{}, {},
{}]".format(i//num_iters_per_epoch, t_loss, t_acc, v_acc, test_acc))
                            temp_output.append([t_loss, t_acc, v_acc, test_acc])
                    output_data[h].append(temp_output)

    np.save('Q1-2-1.npy', output_data)
    return output_data

def number_of_layers():
    # Constants
    B = 250
    iters = 5000
    learning_rates = [0.01, 0.005, 0.001, 0.0005, 0.0001]
    hidden_units = [500]
    output_data = [[]]

    # Load data
    (trainData, trainTarget, validData, validTarget,
        testData, testTarget) = load_notMNIST()

    # Precalculations
    num_iters_per_epoch = len(trainData)//B # number of iterations we have to do for one epoch
    print("Num epochs = ",iters/num_iters_per_epoch)
    inds = np.arange(trainData.shape[0])

    # Set place-holders & variables
    X = tf.placeholder(tf.float32, shape=(None, trainData.shape[-1]), name='X')
    Y = tf.placeholder(tf.float32, shape=(None, 10), name='Y')
    learning_rate = tf.placeholder(tf.float32, name='learning-rate')

    for h in range(0, len(hidden_units)):
        for lr in range(len(learning_rates)):
            # Build graph
            with tf.variable_scope("layer1_"+str(hidden_units[h])+"_"+str(lr), reuse=tf.AUTO_REUSE):
                s_1 = create_new_layer(X, hidden_units[h])
            x_1 = tf.nn.relu(s_1)
            with tf.variable_scope("layer2_"+str(hidden_units[h])+"_"+str(lr), reuse=tf.AUTO_REUSE):
                s_2 = create_new_layer(x_1, hidden_units[h])
            x_2 = tf.nn.softmax(s_2)
            with tf.variable_scope("layer3_"+str(hidden_units[h])+"_"+str(lr), reuse=tf.AUTO_REUSE):
                s_3 = create_new_layer(x_2, 10)
            x_3 = tf.nn.softmax(s_3)

            # Calculate loss & accuracy
            loss = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=s_3, labels=Y))
            accuracy = tf.reduce_mean(tf.cast(tf.equal(tf.argmax(x_3, 1), tf.argmax(Y, 1)), tf.float32))

            print("Number of hidden layers: 2, Number of hidden units", hidden_units[h])

            with tf.Session() as sess:
```

```
                    with tf.variable_scope("default", reuse=tf.AUTO_REUSE):
                        optimizer = tf.train.AdamOptimizer(learning_rate).minimize(loss)
                        coord = tf.train.Coordinator()
                        threads = tf.train.start_queue_runners(sess=sess, coord=coord)
                        sess.run(tf.global_variables_initializer())
                        sess.run(tf.local_variables_initializer())
                        print("Learning rate = ",learning_rates[lr])
                        temp_output = []
                        for i in range(iters):
                            if (i % num_iters_per_epoch == 0):
                                np.random.shuffle(inds)
                            sess.run([optimizer], feed_dict={learning_rate: learning_rates[lr],
                                    X: trainData[inds[B*(i%num_iters_per_epoch):B*((i+1)%num_iters_per_epoch)]],
                                    Y: trainTarget[inds[B*(i%num_iters_per_epoch):B*((i+1)%num_iters_per_epoch)]]})
                            if (i % num_iters_per_epoch == 0):
                                t_loss, t_acc = sess.run([loss, accuracy], feed_dict={X: trainData, Y: trainTarget})
                                v_loss, v_acc = sess.run([loss, accuracy], feed_dict={X: validData, Y: validTarget})
                                test_loss, test_acc = sess.run([loss, accuracy], feed_dict={X: testData, Y: testTarget})
                                print("Epoch: {}, Training Loss: {}, Accuracies: [{}, {},
{}]".format(i//num_iters_per_epoch, t_loss, t_acc, v_acc, test_acc))
                                temp_output.append([t_loss, t_acc, v_acc, test_acc])
                        output_data[h].append(temp_output)

    np.save('Q1-2-2.npy', output_data)
    return output_data

#output = number_of_hidden_units()
output = number_of_layers()
```

regularization.py

```python
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt

def dropout(x, is_training, p):
    return tf.cond(is_training, lambda: tf.nn.dropout(x, p, name='dropout'), lambda: tf.identity(x))

# Data loader for notMNIST dataset
def load_notmnist_data():
    with np.load("notMNIST.npz") as data:
        Data, Target = data["images"], data["labels"]
        np.random.seed(521)
        randIndx = np.arange(len(Data))
        np.random.shuffle(randIndx)
        Data = Data[randIndx]/255
        Target = Target[randIndx]
        trainData, trainTarget = Data[:15000], Target[:15000]
        t = np.zeros((trainTarget.shape[0], 10))
        t[np.arange(trainTarget.shape[0]), trainTarget] = 1
        trainTarget = t
        validData, validTarget = Data[15000:16000], Target[15000:16000]
        t = np.zeros((validTarget.shape[0], 10))
        t[np.arange(validTarget.shape[0]), validTarget] = 1
        validTarget = t
        testData, testTarget = Data[16000:], Target[16000:]
        t = np.zeros((testTarget.shape[0], 10))
        t[np.arange(testTarget.shape[0]), testTarget] = 1
        testTarget = t
        return (trainData.reshape(trainData.shape[0], -1), trainTarget, validData.reshape(validData.shape[0], -1),
validTarget, testData.reshape(testData.shape[0], -1), testTarget)


def FCN(x, depth, name, use_dropout=False, is_training=tf.constant(False), use_relu=False):
    W = tf.get_variable(name=name + "_W", shape=(x.shape[1], depth), dtype=tf.float64)
    b = tf.get_variable(name=name+ "_b", shape=(depth,), dtype=tf.float64, initializer=tf.zeros_initializer)
    if use_dropout:
        if use_relu:
            return dropout(tf.nn.relu(tf.matmul(x, W) + b), is_training, 0.5)
        else:
            return dropout(tf.matmul(x, W) + b, is_training, 0.5)
    else:
        if use_relu:
            return tf.nn.relu(tf.matmul(x, W) + b)
        else:
            tf.matmul(x, W) + b

def build_network(input_node, is_training_t):
```

```python
    #can be changed for 2-layer networks
    num_hidden_units = 1000
    L1_out = FCN(input_node[0], num_hidden_units, name='Layer_1', use_dropout=True, is_training=is_training_t,
             use_relu=True)
    W = tf.get_variable(name="output_W", shape=(L1_out.shape[1], 10), dtype=tf.float64)
    b = tf.get_variable(name="output_b", shape=(10,), dtype=tf.float64, initializer=tf.zeros_initializer)

    #for multiple nerual networks
    # L1_out = FCN(input_node[0], num_hidden_units, name='Layer_1', use_dropout=True, is_training=is_training_t,
use_relu=True)
    # L2_out = FCN(L1_out, num_hidden_units, name='Layer_2', use_dropout=True, is_training=is_training_t,
    #            use_relu=True)
    # L3_out = FCN(L2_out, num_hidden_units, name='Layer_3', use_dropout=True, is_training=is_training_t,
    #            use_relu=True)
    # L4_out = FCN(L3_out, num_hidden_units, name='Layer_4', use_dropout=True, is_training=is_training_t,
    #            use_relu=True)
    # W = tf.get_variable(name="output_W", shape=(L1_out.shape[1], 10), dtype=tf.float64)
    # b = tf.get_variable(name="output_b", shape=(10,), dtype=tf.float64, initializer=tf.zeros_initializer)

    # y_pred_raw = tf.matmul(L4_out, W) + b

    y_pred_raw = tf.matmul(L1_out, W) + b
    return y_pred_raw

def learning():
    xTrain, yTrain, xValid, yValid, xTest, yTest = load_notmnist_data()

    with tf.Graph().as_default():
        num_hidden_units = 1000
        decay = 0
        B = 500
        learning_rates = [0.01, 0.005, 0.001]
        iters = 5000
        max_num_epochs = (B*iters)//len(xTrain)
        if B*iters % len(xTrain):
            max_num_epochs += 1
        num_iters_per_epoch = len(xTrain) // B  # number of iterations we have to do for one epoch
        print("Num epochs = ", iters / num_iters_per_epoch)

        # hyperparameters
        learning_rate = tf.placeholder(dtype=tf.float64, name="learning-rate")
        is_training_t = tf.placeholder(dtype=tf.bool, name="is_training")

        base_iterator = tf.data.Iterator.from_structure((tf.float64, tf.float64), ((None, 784), (None, 10)))
        input_node = base_iterator.get_next()
        y_pred_raw = build_network(input_node, is_training_t)

        y_pred = tf.nn.softmax(y_pred_raw)
        CE_loss = tf.losses.softmax_cross_entropy(input_node[1], y_pred_raw)

        vars = tf.global_variables()
        l2s = []
        for var in vars:
            l2s.append(tf.nn.l2_loss(var))
        l2_loss = tf.reduce_sum(tf.stack(l2s, axis=0))
        total_loss = CE_loss + decay * l2_loss
        accuracy = tf.reduce_mean(tf.cast(tf.equal(tf.argmax(y_pred, 1), tf.argmax(input_node[1], 1)), tf.float32))

        # optimizer function
        optimizer = tf.train.AdamOptimizer(learning_rate).minimize(total_loss)

        X = tf.placeholder(dtype=tf.float64, name="X")
        Y = tf.placeholder(dtype=tf.float64, name="Y")
        Xdata = tf.data.Dataset.from_tensor_slices(X)
        Ydata = tf.data.Dataset.from_tensor_slices(Y)
        sample_dataset = tf.data.Dataset.zip((Xdata, Ydata))
        batched_dataset = sample_dataset.batch(B)
        # TODO run a lot of iterations, plot loss vs epochs and classification error vs epochs
        accuracy_list = []
        ce_list = []
        check_points = [iters//4, iters//2, 3*iters//4, iters-1]
        saver = tf.train.Saver(vars)
        with tf.Session() as sess:
            sess.run(tf.global_variables_initializer())
            # initialize data input pippeline for training
            dataset_init = base_iterator.make_initializer(batched_dataset)
```

```
            for i in range(max_num_epochs):
                sess.run(dataset_init, feed_dict={X:xTrain, Y:yTrain})
                j = 0
                while True:
                    try:
                        sess.run([optimizer, ], feed_dict={learning_rate: 0.005, is_training_t: True})
                        j += 1
                        if i * num_iters_per_epoch + j in check_points:
                            saver.save(sess, '.\my_model', global_step=i)
                    except tf.errors.OutOfRangeError:
                        break

                # initialize data iterator for getting numbers to plot
                # on train
                sess.run(dataset_init, feed_dict={X: xTrain, Y: yTrain})
                this_acc = 0.0
                this_ce = 0.0
                j = 0
                while True:
                    try:
                        acc, ce = sess.run([accuracy, CE_loss], feed_dict={is_training_t: False})
                        this_acc += acc
                        this_ce += ce
                        j += 1
                    except tf.errors.OutOfRangeError:
                        break
                train_acc = this_acc/j
                train_ce = this_ce/j
                # on val
                sess.run(dataset_init, feed_dict={X: xValid, Y: yValid})
                this_acc = 0.0
                this_ce = 0.0
                j = 0
                while True:
                    try:
                        acc, ce = sess.run([accuracy, CE_loss], feed_dict={is_training_t: False})
                        this_acc += acc
                        this_ce += ce
                        j += 1
                    except tf.errors.OutOfRangeError:
                        break
                val_acc = this_acc / j
                val_ce = this_ce / j
                # on test
                sess.run(dataset_init, feed_dict={X: xTest, Y: yTest})
                this_acc = 0.0
                this_ce = 0.0
                j = 0
                while True:
                    try:
                        acc, ce = sess.run([accuracy, CE_loss], feed_dict={is_training_t: False})
                        this_acc += acc
                        this_ce += ce
                        j += 1
                    except tf.errors.OutOfRangeError:
                        break
                test_acc = this_acc / j
                test_ce = this_ce / j
                accuracy_list.append((train_acc, val_acc, test_acc))
                ce_list.append((train_ce, val_ce, test_ce))
                print("Epoch: {}, Training Loss: {}, Accuracies: [{}, {}, {}]".format(i,
                                                            train_ce, train_acc, val_acc,
                                                            test_acc))
    return accuracy_list, ce_list


def visualization(filepath, index=1):
    base_iterator = tf.data.Iterator.from_structure((tf.float64, tf.float64), ((None, 784), (None, 10)))
    input_node = base_iterator.get_next()
    is_training_t = tf.placeholder(dtype=tf.bool, name="is_training")
    _ = build_network(input_node, is_training_t)
    saver = tf.train.Saver(tf.global_variables())
    for var in tf.global_variables():
        if var.name == "Layer_1_W:0":
            l1_w = var
    with tf.Session() as sess:
        saver.restore(sess, filepath)
```

```
        layer1_W = sess.run(l1_w)
        target = layer1_W[:, index]
        plt.imshow(np.reshape(target, (28,28)))
        plt.show()


if __name__ == "__main__":
    #accs, ces = learning()
    #acc_array = np.array(accs)
    #x = np.arange(acc_array.shape[0])
    #plt.plot(x, acc_array)
    #plt.show()
    #ces_array = np.array(ces)
    #plt.plot(x, ces_array)
    #plt.show()

    (ac, ce) = learning()
```

exhaustive.py

```python
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import math
import time

def dropout(x, is_training, p):
    return tf.cond(is_training, lambda: tf.nn.dropout(x, p, name='dropout'), lambda: tf.identity(x))

# Data loader for notMNIST dataset
def load_notmnist_data():
    with np.load("notMNIST.npz") as data:
        Data, Target = data["images"], data["labels"]
        np.random.seed(521)
        randIndx = np.arange(len(Data))
        np.random.shuffle(randIndx)
        Data = Data[randIndx]/255
        Target = Target[randIndx]
        trainData, trainTarget = Data[:15000], Target[:15000]
        t = np.zeros((trainTarget.shape[0], 10))
        t[np.arange(trainTarget.shape[0]), trainTarget] = 1
        trainTarget = t
        validData, validTarget = Data[15000:16000], Target[15000:16000]
        t = np.zeros((validTarget.shape[0], 10))
        t[np.arange(validTarget.shape[0]), validTarget] = 1
        validTarget = t
        testData, testTarget = Data[16000:], Target[16000:]
        t = np.zeros((testTarget.shape[0], 10))
        t[np.arange(testTarget.shape[0]), testTarget] = 1
        testTarget = t
        return (trainData.reshape(trainData.shape[0], -1), trainTarget, validData.reshape(validData.shape[0], -1),
validTarget, testData.reshape(testData.shape[0], -1), testTarget)


def FCN(x, depth, name, use_dropout=False, is_training=tf.constant(False), use_relu=False):
    W = tf.get_variable(name=name + "_W", shape=(x.shape[1], depth), dtype=tf.float64)
    b = tf.get_variable(name=name+ "_b", shape=(depth,), dtype=tf.float64, initializer=tf.zeros_initializer)
    if use_dropout:
        if use_relu:
            return dropout(tf.nn.relu(tf.matmul(x, W) + b), is_training, 0.5)
        else:
            return dropout(tf.matmul(x, W) + b, is_training, 0.5)
    else:
        if use_relu:
            return tf.nn.relu(tf.matmul(x, W) + b)
        else:
            tf.matmul(x, W) + b

def build_network(input_node, is_training_t, dropout, num_hidden_units, num_layers):

    L_out = []

    #can be changed for 2-layer networks
    L_out.append(FCN(input_node[0], num_hidden_units, name='Layer_1', use_dropout=dropout, is_training=is_training_t,
            use_relu=True))

    if num_layers >= 2:
        for i in range(1,num_layers):
```

```python
        L_out.append(FCN(L_out[i-1], num_hidden_units, name='Layer_'+str(i+1), use_dropout=dropout,
is_training=is_training_t,
                    use_relu=True))

    W = tf.get_variable(name="output_W", shape=(L_out[0].shape[1], 10), dtype=tf.float64)
    b = tf.get_variable(name="output_b", shape=(10,), dtype=tf.float64, initializer=tf.zeros_initializer)


    #for multiple nerual networks
    # L1_out = FCN(input_node[0], num_hidden_units, name='Layer_1', use_dropout=True, is_training=is_training_t,
use_relu=True)
    # L2_out = FCN(L1_out, num_hidden_units, name='Layer_2', use_dropout=True, is_training=is_training_t,
    #             use_relu=True)
    # L3_out = FCN(L2_out, num_hidden_units, name='Layer_3', use_dropout=True, is_training=is_training_t,
    #             use_relu=True)
    # L4_out = FCN(L3_out, num_hidden_units, name='Layer_4', use_dropout=True, is_training=is_training_t,
    #             use_relu=True)
    # W = tf.get_variable(name="output_W", shape=(L1_out.shape[1], 10), dtype=tf.float64)
    # b = tf.get_variable(name="output_b", shape=(10,), dtype=tf.float64, initializer=tf.zeros_initializer)

    # y_pred_raw = tf.matmul(L4_out, W) + b

    y_pred_raw = tf.matmul(L_out[len(L_out)-1], W) + b
    return y_pred_raw

def learning():
    xTrain, yTrain, xValid, yValid, xTest, yTest = load_notmnist_data()

    with tf.Graph().as_default():

        np.random.seed(int(time.time()))
        num_hidden_units = np.random.randint(100, 500 + 1)
        decay = math.exp((-6 + -9) * np.random.random_sample() - 9)
        B = 500
        dropout = np.random.choice([True, False])
        num_layers = np.random.randint(1, 5 + 1)
        lr = math.exp((-3.5 + 7.5) * np.random.random_sample() - 7.5)
        iters = 5000
        max_num_epochs = (B*iters)//len(xTrain)
        if B*iters % len(xTrain):
            max_num_epochs += 1
        num_iters_per_epoch = len(xTrain) // B  # number of iterations we have to do for one epoch
        print("Number of hidden units", num_hidden_units)
        print("Decay", decay)
        print("Dropout", dropout)
        print("Num layers", num_layers)
        print("Learning rate", lr)
        print("Num epochs = ", iters / num_iters_per_epoch)

        # hyperparameters
        learning_rate = tf.placeholder(dtype=tf.float64, name="learning-rate")
        is_training_t = tf.placeholder(dtype=tf.bool, name="is_training")

        base_iterator = tf.data.Iterator.from_structure((tf.float64, tf.float64), ((None, 784), (None, 10)))
        input_node = base_iterator.get_next()
        y_pred_raw = build_network(input_node, is_training_t, dropout, num_hidden_units, num_layers)

        y_pred = tf.nn.softmax(y_pred_raw)
        CE_loss = tf.losses.softmax_cross_entropy(input_node[1], y_pred_raw)

        vars = tf.global_variables()
        l2s = []
        for var in vars:
            l2s.append(tf.nn.l2_loss(var))
        l2_loss = tf.reduce_sum(tf.stack(l2s, axis=0))
        total_loss = CE_loss + decay * l2_loss
        accuracy = tf.reduce_mean(tf.cast(tf.equal(tf.argmax(y_pred, 1), tf.argmax(input_node[1], 1)), tf.float32))

        # optimizer function
        optimizer = tf.train.AdamOptimizer(learning_rate).minimize(total_loss)

        X = tf.placeholder(dtype=tf.float64, name="X")
        Y = tf.placeholder(dtype=tf.float64, name="Y")
        Xdata = tf.data.Dataset.from_tensor_slices(X)
        Ydata = tf.data.Dataset.from_tensor_slices(Y)
        sample_dataset = tf.data.Dataset.zip((Xdata, Ydata))
        batched_dataset = sample_dataset.batch(B)
```

```
        # TODO run a lot of iterations, plot loss vs epochs and classification error vs epochs
        accuracy_list = []
        ce_list = []
        check_points = [iters//4, iters//2, 3*iters//4, iters-1]
        saver = tf.train.Saver(vars)
        with tf.Session() as sess:
            sess.run(tf.global_variables_initializer())
            # initialize data input pippeline for training
            dataset_init = base_iterator.make_initializer(batched_dataset)
            for i in range(max_num_epochs):
                sess.run(dataset_init, feed_dict={X:xTrain, Y:yTrain})
                j = 0
                while True:
                    try:
                        sess.run([optimizer, ], feed_dict={learning_rate: lr, is_training_t: True})
                        j += 1
                        if i * num_iters_per_epoch + j in check_points:
                            saver.save(sess, '.\my_model', global_step=i)
                    except tf.errors.OutOfRangeError:
                        break

                # initialize data iterator for getting numbers to plot
                # on train
                sess.run(dataset_init, feed_dict={X: xTrain, Y: yTrain})
                this_acc = 0.0
                this_ce = 0.0
                j = 0
                while True:
                    try:
                        acc, ce = sess.run([accuracy, CE_loss], feed_dict={is_training_t: False})
                        this_acc += acc
                        this_ce += ce
                        j += 1
                    except tf.errors.OutOfRangeError:
                        break
                train_acc = this_acc/j
                train_ce = this_ce/j
                # on val
                sess.run(dataset_init, feed_dict={X: xValid, Y: yValid})
                this_acc = 0.0
                this_ce = 0.0
                j = 0
                while True:
                    try:
                        acc, ce = sess.run([accuracy, CE_loss], feed_dict={is_training_t: False})
                        this_acc += acc
                        this_ce += ce
                        j += 1
                    except tf.errors.OutOfRangeError:
                        break
                val_acc = this_acc / j
                val_ce = this_ce / j
                # on test
                sess.run(dataset_init, feed_dict={X: xTest, Y: yTest})
                this_acc = 0.0
                this_ce = 0.0
                j = 0
                while True:
                    try:
                        acc, ce = sess.run([accuracy, CE_loss], feed_dict={is_training_t: False})
                        this_acc += acc
                        this_ce += ce
                        j += 1
                    except tf.errors.OutOfRangeError:
                        break
                test_acc = this_acc / j
                test_ce = this_ce / j
                accuracy_list.append((train_acc, val_acc, test_acc))
                ce_list.append((train_ce, val_ce, test_ce))
                print("Epoch: {}, Training Loss: {}, Accuracies: [{}, {}, {}]".format(i,
                                                            train_ce, train_acc, val_acc,
                                                            test_acc))

    return accuracy_list, ce_list


def visualization(filepath, index=1):
    base_iterator = tf.data.Iterator.from_structure((tf.float64, tf.float64), ((None, 784), (None, 10)))
```

```
    input_node = base_iterator.get_next()
    is_training_t = tf.placeholder(dtype=tf.bool, name="is_training")
    _ = build_network(input_node, is_training_t)
    saver = tf.train.Saver(tf.global_variables())
    for var in tf.global_variables():
        if var.name == "Layer_1_W:0":
            l1_w = var
    with tf.Session() as sess:
        saver.restore(sess, filepath)
        layer1_W = sess.run(l1_w)
        target = layer1_W[:, index]
        plt.imshow(np.reshape(target, (28,28)))
        plt.show()


if __name__ == "__main__":
    #accs, ces = learning()
    #acc_array = np.array(accs)
    #x = np.arange(acc_array.shape[0])
    #plt.plot(x, acc_array)
    #plt.show()
    #ces_array = np.array(ces)
    #plt.plot(x, ces_array)
    #plt.show()

    (ac, ce) = learning()
```